

# gwp1\_grp\_8781\_Assignment\_Notebook

April 15, 2025

## Group Members

1. Halkano Molu Guracha
2. Biswajit Paljit
3. Joel Muchemeranwa

### 0.0.1 Installing Dependancies

```
[1]: !pip install nelson_siegel_svensson  
!pip install yfinance
```

Collecting nelson\_siegel\_svensson

Downloading nelson\_siegel\_svensson-0.5.0-py2.py3-none-any.whl.metadata (6.7 kB)

Requirement already satisfied: Click>=8.0 in /usr/local/lib/python3.11/dist-packages (from nelson\_siegel\_svensson) (8.1.8)

Requirement already satisfied: numpy>=1.22 in /usr/local/lib/python3.11/dist-packages (from nelson\_siegel\_svensson) (2.0.2)

Requirement already satisfied: scipy>=1.7 in /usr/local/lib/python3.11/dist-packages (from nelson\_siegel\_svensson) (1.14.1)

Requirement already satisfied: matplotlib>=3.5 in /usr/local/lib/python3.11/dist-packages (from nelson\_siegel\_svensson) (3.10.0)

Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib>=3.5->nelson\_siegel\_svensson) (1.3.1)

Requirement already satisfied: cyclor>=0.10 in /usr/local/lib/python3.11/dist-packages (from matplotlib>=3.5->nelson\_siegel\_svensson) (0.12.1)

Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib>=3.5->nelson\_siegel\_svensson) (4.57.0)

Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib>=3.5->nelson\_siegel\_svensson) (1.4.8)

Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib>=3.5->nelson\_siegel\_svensson) (24.2)

Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.11/dist-packages (from matplotlib>=3.5->nelson\_siegel\_svensson) (11.1.0)

Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib>=3.5->nelson\_siegel\_svensson) (3.2.3)

Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.11/dist-packages (from matplotlib>=3.5->nelson\_siegel\_svensson) (2.8.2)

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-

packages (from python-dateutil>=2.7->matplotlib>=3.5->nelson\_siegel\_svensson)  
 (1.17.0)  
 Downloading nelson\_siegel\_svensson-0.5.0-py2.py3-none-any.whl (9.9 kB)  
 Installing collected packages: nelson\_siegel\_svensson  
 Successfully installed nelson\_siegel\_svensson-0.5.0  
 Requirement already satisfied: yfinance in /usr/local/lib/python3.11/dist-  
 packages (0.2.55)  
 Requirement already satisfied: pandas>=1.3.0 in /usr/local/lib/python3.11/dist-  
 packages (from yfinance) (2.2.2)  
 Requirement already satisfied: numpy>=1.16.5 in /usr/local/lib/python3.11/dist-  
 packages (from yfinance) (2.0.2)  
 Requirement already satisfied: requests>=2.31 in /usr/local/lib/python3.11/dist-  
 packages (from yfinance) (2.32.3)  
 Requirement already satisfied: multitasking>=0.0.7 in  
 /usr/local/lib/python3.11/dist-packages (from yfinance) (0.0.11)  
 Requirement already satisfied: platformdirs>=2.0.0 in  
 /usr/local/lib/python3.11/dist-packages (from yfinance) (4.3.7)  
 Requirement already satisfied: pytz>=2022.5 in /usr/local/lib/python3.11/dist-  
 packages (from yfinance) (2025.2)  
 Requirement already satisfied: frozendict>=2.3.4 in  
 /usr/local/lib/python3.11/dist-packages (from yfinance) (2.4.6)  
 Requirement already satisfied: peewee>=3.16.2 in /usr/local/lib/python3.11/dist-  
 packages (from yfinance) (3.17.9)  
 Requirement already satisfied: beautifulsoup4>=4.11.1 in  
 /usr/local/lib/python3.11/dist-packages (from yfinance) (4.13.3)  
 Requirement already satisfied: soupsieve>1.2 in /usr/local/lib/python3.11/dist-  
 packages (from beautifulsoup4>=4.11.1->yfinance) (2.6)  
 Requirement already satisfied: typing-extensions>=4.0.0 in  
 /usr/local/lib/python3.11/dist-packages (from beautifulsoup4>=4.11.1->yfinance)  
 (4.13.1)  
 Requirement already satisfied: python-dateutil>=2.8.2 in  
 /usr/local/lib/python3.11/dist-packages (from pandas>=1.3.0->yfinance) (2.8.2)  
 Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-  
 packages (from pandas>=1.3.0->yfinance) (2025.2)  
 Requirement already satisfied: charset-normalizer<4,>=2 in  
 /usr/local/lib/python3.11/dist-packages (from requests>=2.31->yfinance) (3.4.1)  
 Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-  
 packages (from requests>=2.31->yfinance) (3.10)  
 Requirement already satisfied: urllib3<3,>=1.21.1 in  
 /usr/local/lib/python3.11/dist-packages (from requests>=2.31->yfinance) (2.3.0)  
 Requirement already satisfied: certifi>=2017.4.17 in  
 /usr/local/lib/python3.11/dist-packages (from requests>=2.31->yfinance)  
 (2025.1.31)  
 Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-  
 packages (from python-dateutil>=2.8.2->pandas>=1.3.0->yfinance) (1.17.0)

```
[2]: import pandas as pd
import matplotlib.pyplot as plt
from scipy.interpolate import CubicSpline
from nelson_siegel_svensson.calibrate import calibrate_ns_ols
import numpy as np
from numpy import linalg as LA
import seaborn as sns
```

## 1 Task 2

### 1.0.1 Q2. a. and Q2. b.

We selected the country - India. The following dataset contains the daily yields of the Indian bonds with the following maturities from 2006 to 2025(current).

| Maturity | Source                                   |
|----------|--|
| 3 Months | <a href="#">India 3-Month Bond Yield</a> |
| 6 Months | <a href="#">India 6-Month Bond Yield</a> |
| 1 Year   | <a href="#">India 1-Year Bond Yield</a>  |
| 2 Year   | <a href="#">India 2-Year Bond Yield</a>  |
| 3 Year   | <a href="#">India 3-Year Bond Yield</a>  |
| 4 Year   | <a href="#">India 4-Year Bond Yield</a>  |
| 5 Year   | <a href="#">India 5-Year Bond Yield</a>  |
| 6 Year   | <a href="#">India 6-Year Bond Yield</a>  |
| 7 Year   | <a href="#">India 7-Year Bond Yield</a>  |
| 8 Year   | <a href="#">India 8-Year Bond Yield</a>  |
| 9 Year   | <a href="#">India 9-Year Bond Yield</a>  |
| 10 Year  | <a href="#">India 10-Year Bond Yield</a> |
| 12 Year  | <a href="#">India 12-Year Bond Yield</a> |
| 15 Year  | <a href="#">India 15-Year Bond Yield</a> |
| 24 Year  | <a href="#">India 24-Year Bond Yield</a> |
| 30 Year  | <a href="#">India 30-Year Bond Yield</a> |

```
[5]: df = pd.read_csv('C:/Users/Gast01/WQU-GWP1/India Bond Yield Data.csv') # Change
      ↪the file path as the path in your device.
```

```
df.index = pd.to_datetime(df['Date'])
```

```
df = df.drop('Date', axis=1)
```

```
df = df.dropna()
```

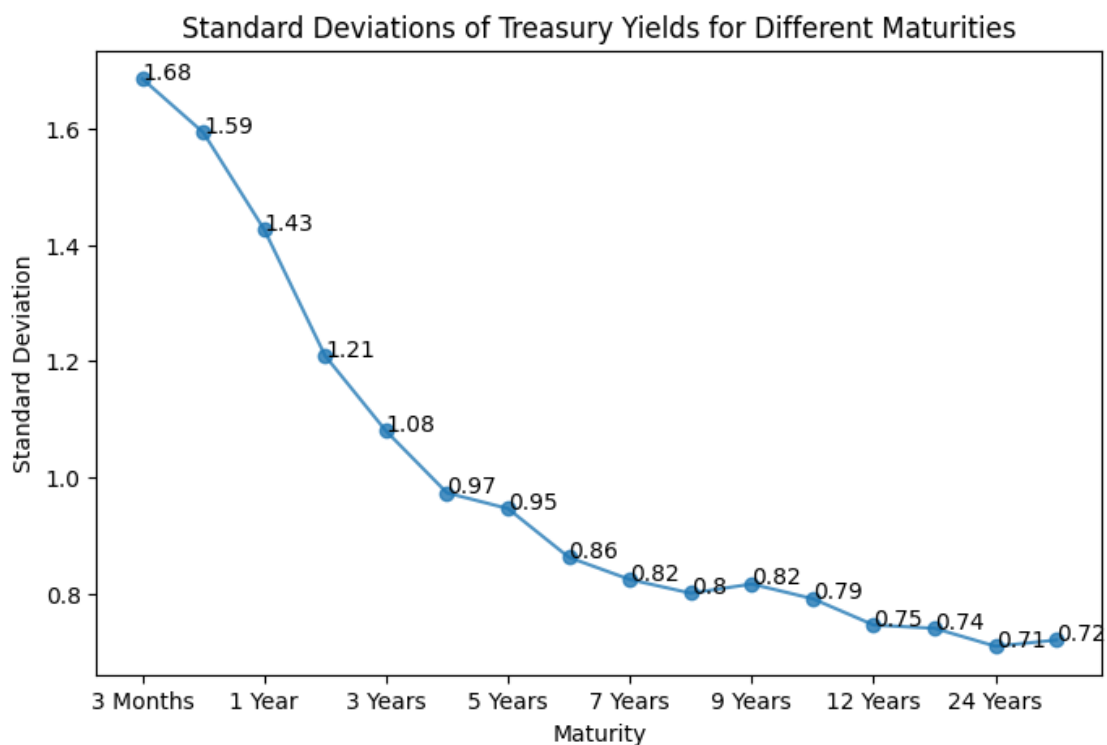
```
[6]: def plot_std(df):
      """
      Plotting the standard deviation of the treasury yields for different
      ↪maturities
      """
```

```

y_std = df.std()
fig, ax = plt.subplots()
y_std.plot(figsize = (8,5),marker='o', title='Standard Deviations of Treasury_
↳Yields for Different Maturities', alpha=0.8)
plt.xlabel("Maturity")
plt.ylabel("Standard Deviation")
for i in range(len(y_std)):
    ax.annotate(str(round(y_std.iloc[i],2)),xy=(i,y_std.iloc[i]))
plt.show()

```

```
[7]: plot_std(df)
```



```

[8]: def plot_yield_curve(date, fig_n):
    """
    Plotting the yield curve for a given date
    """

    maturities = df.columns # Maturities
    fig, ax = plt.subplots(figsize=(6.15, 4))
    ax.plot(maturities, df.loc[date], marker='D', label='Yield Curve at ' +
↳date)

```

```

ax.set_yticklabels(['{:0.2f}%'.format(y) for y in ax.get_yticks()])
ax.set_xticks(range(len(maturities)))
labels = [m if i % 5 == 0 else '' for i, m in enumerate(maturities)]
ax.set_xticklabels(labels)

# Add labels and title
ax.set_xlabel('Maturity')
ax.set_ylabel('Yield')
ax.set_title(fig_n+f'Treasury Yield Curve as of {date}')

# Show the plot
plt.grid(False)
plt.show()

```

### 1.0.2 Q2. c.

Fitting a Nelson Siegel Model on all the maturities on a specific date and comparing it with the original yield curve as of that date

```

[9]: def fit_ns(t, y, tau0=1.0):
    """
    Fitting the Nelson-Siegel model to the yield curve data for a given date.
    """

    curve, status = calibrate_ns_ols(t, y, tau0=1.0) # starting value of 1.0 for
    ↪the optimization of tau
    assert status.success

    print(curve)

    return curve

```

```

[10]: def plot_ns(date, y_hat, t_hat):
    """
    Plotting the yield curve for a given date.
    """

    plt.plot(t_hat, y_hat(t_hat))
    plt.xlabel("Maturity")
    plt.ylabel("Yield")
    plt.title(f"NS Model Result as of {date}")

    plt.show()

```

```

[11]: date = "2020-01-10"
t = np.array([0.25,0.5,1,2,3,4,5,6,7,8,9,10,12,15,24,30])

```

```

y = np.array(df.loc[date])
curve = fit_ns(t, y, tau0=1.0)
y_hat = curve
t_hat = np.linspace(0.5,30,100)
plot_yield_curve(date, '')
plot_ns(date, y_hat, t_hat)

```

```

NelsonSiegelCurve(beta0=np.float64(7.262536953580489),
beta1=np.float64(-2.2736710165571905), beta2=np.float64(-1.7792170270160836),
tau=np.float64(0.9062504122308159))

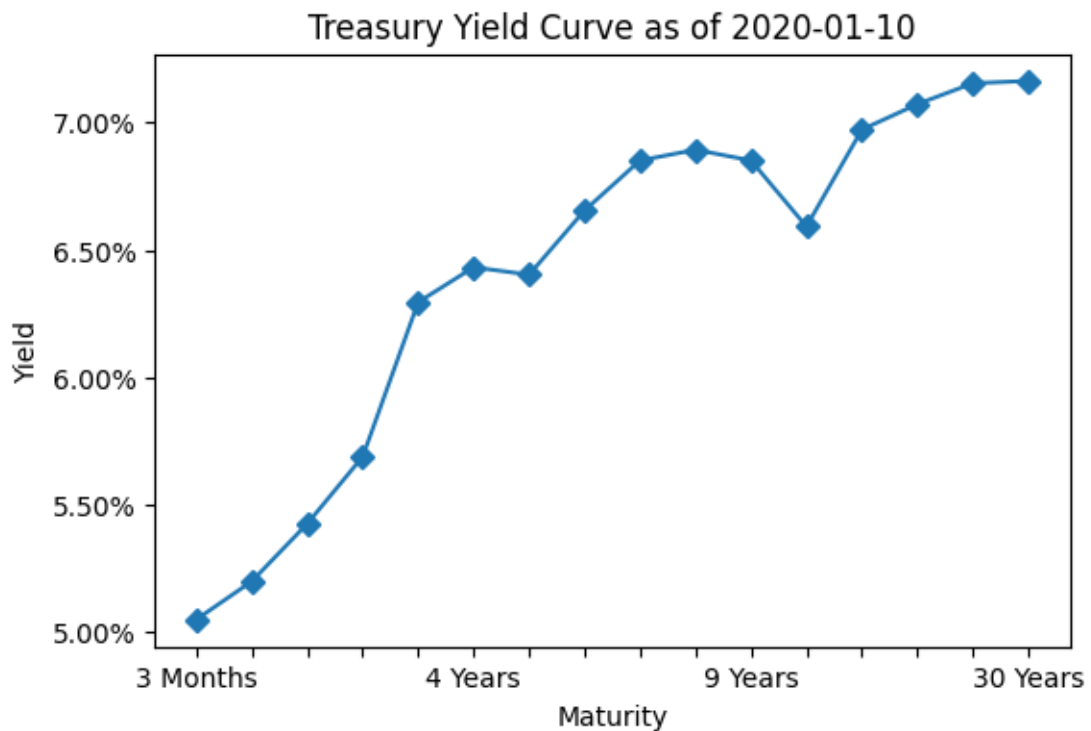
```

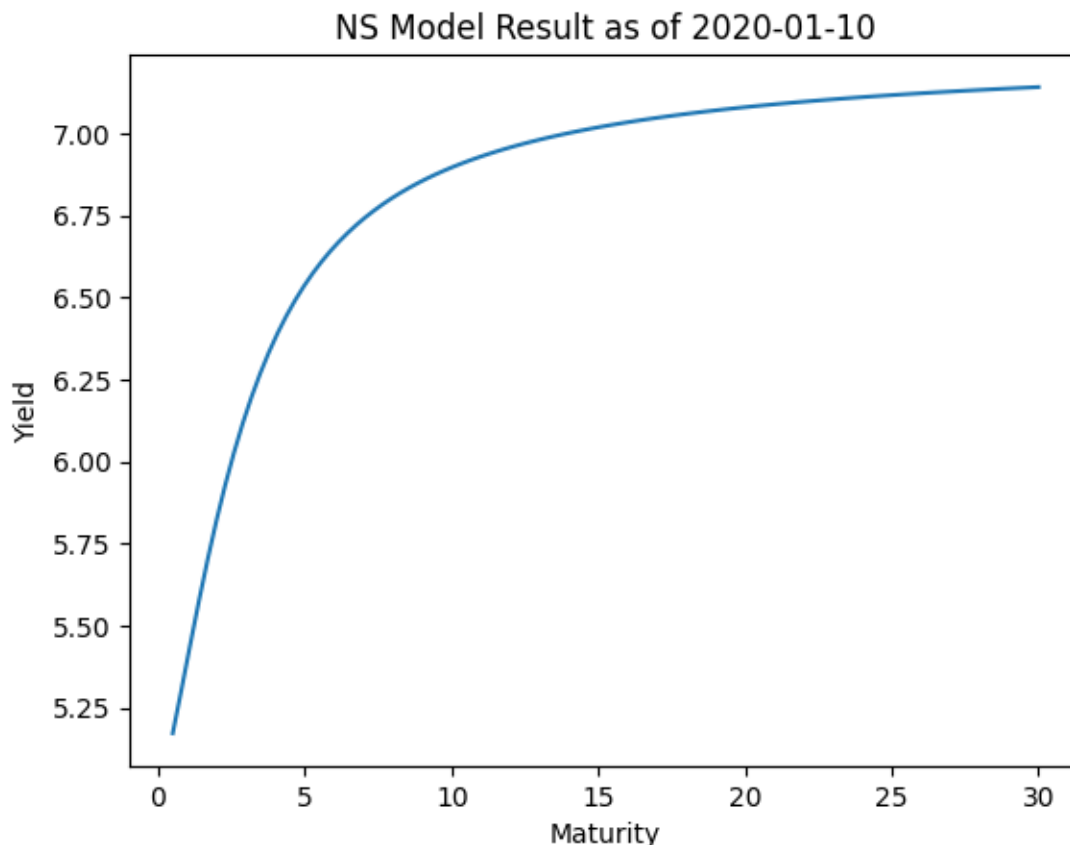
<ipython-input-8-f6116200d263>:10: UserWarning: set\_ticklabels() should only be used with a fixed number of ticks, i.e. after set\_ticks() or using a FixedLocator.

```

ax.set_yticklabels(['{:.2f}%'.format(y) for y in ax.get_yticks()])

```





### 1.0.3 Q2. f.

In the solution above, we have fit a Nelson Siegel Model on the yield curve as of 2020-01-10. The result `NelsonSiegelCurve(beta0=np.float64(7.262536960509667), beta1=np.float64(-2.273671047246356), beta2=np.float64(-1.779216868209846), tau=np.float64(0.9062504579008243))` indicates that the  $\beta_0 = 7.26$  showing the level of the yield curve.  $\beta_1 = -2.27$  shows the slope of the yield curve and the  $\beta_2 = -1.77$  shows the shape or the curvature of the yield curve. the decay rate indicated by  $\tau = 0.91$  showing a slow rate of decay.

A level of 7.26 shows that the long term expectation of the yields to be around 7%.

A negative slope can be attributed to the fact that there are dips in the long term yields, for example the 10 year yield is smaller than the 6 years, 7years and 8 years yields.

A negative curvature indicates a concave behaviour. The yields started rising steep in the short run, but flatten over a longer period.

This interpretation is in line with the mid pandemic situation in 2020. Complete economic shut-downs has raised riskiness of the short term borrowings, hence the rise in yields sharply.

A smaller value of  $\tau$  (close to 1) indicates that the effects of the slope and curvature parameters decay relatively slowly as maturity increases. This suggests that the short-to-medium-term rates

have a strong influence on the shape of the curve, but as maturity increases, the curve flattens, and the level dominates.

#### 1.0.4 Q2. d. and Q2. f.

Fitting a Cubic Spline Model on all the maturities on a specific date and comparing it with the original yield curve as of that date.

Since there are 16 maturities, there would be 15 splines that pass through the 16 points. The splines can be shown as follows:

$$f(x) = a_1x^3 + b_1x^2 + c_1x + d_1, \quad \text{when } 0.25 \leq x \leq 0.5$$

$$f(x) = a_2x^3 + b_2x^2 + c_2x + d_2, \quad \text{when } 0.5 \leq x \leq 1$$

$$f(x) = a_3x^3 + b_3x^2 + c_3x + d_3, \quad \text{when } 1 \leq x \leq 2$$

and so on...

From the above equations, we have (  $15 \times 4 = 60$  ) unknowns. Hence, we need 60 equations to solve for the parameters.

Thus plugging each boundary, we get 30 equations as shown below:

$$a_1(0.25)^3 + b_1(0.25)^2 + c_1(0.25) + d_1 = 5.05 \quad (1)$$

$$a_1(0.5)^3 + b_1(0.5)^2 + c_1(0.5) + d_1 = 5.2 \quad (2)$$

$$a_2(0.5)^3 + b_2(0.5)^2 + c_2(0.5) + d_2 = 5.2 \quad (3)$$

$$a_2(1)^3 + b_2(1)^2 + c_2(1) + d_2 = 5.42 \quad (4)$$

$$a_3(1)^3 + b_3(1)^2 + c_3(1) + d_3 = 5.42 \quad (5)$$

$$a_3(2)^3 + b_3(2)^2 + c_3(2) + d_3 = 5.69 \quad (6)$$

Now since each interior point is a part of 2 splines, their slopes and curvatures must be the same. Therefore, their first order derivatives and their second order derivatives should be the same. Since there are (  $16 - 2 = 14$  ) interior points, we get 14 first order equations, and 14 second order equations.

The first order equations are as follows:

$$3a_1(0.5)^2 + 2b_1(0.5) + c_1 = 3a_2(0.5)^2 + 2b_2(0.5) + c_2 \quad (31)$$

$$3a_2(1)^2 + 2b_2(1) + c_2 = 3a_3(1)^2 + 2b_3(1) + c_3 \quad (32)$$

The second order equations are as follows:

$$6a_1(0.5) + 2b_1 = 6a_2(0.5) + 2b_2 \quad (45)$$



$$6a_2(1) + 2b_2 = 6a_3(1) + 2b_3 \quad (46)$$

In total now we have 58 equations. Finally we assume the “Natural End Condition” and consider the second order derivatives of the exterior points to be 0/ Since there are 2 exterior points, we get 2 equations, which gives us full 60 equations.

$$6a_1(0.25) + 2b_1 = 0 \quad (59)$$

$$6a_3(30) + 2b_3 = 0 \quad (60)$$

To calculate the cubic spline model, we use the `CubicSpline` function from the **SciPy** library, typically imported as `cs`.

The command is:

```
cs = CubicSpline(t, y, bc_type=((2, 0.0), (2, 0.0)))
```

where:

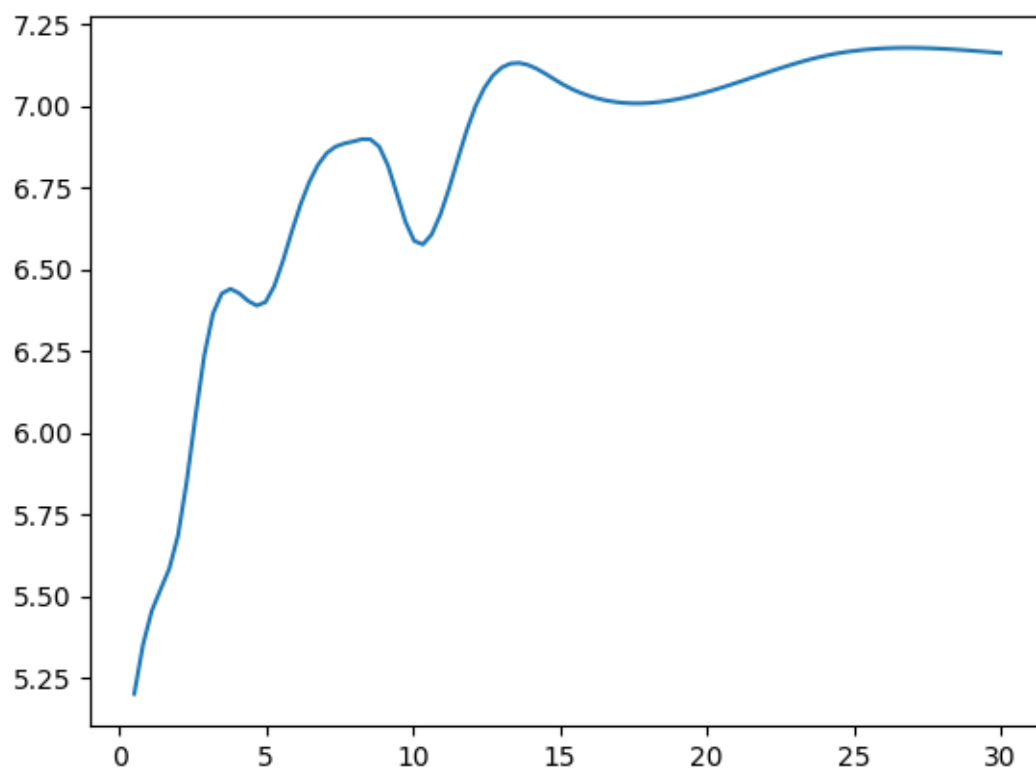
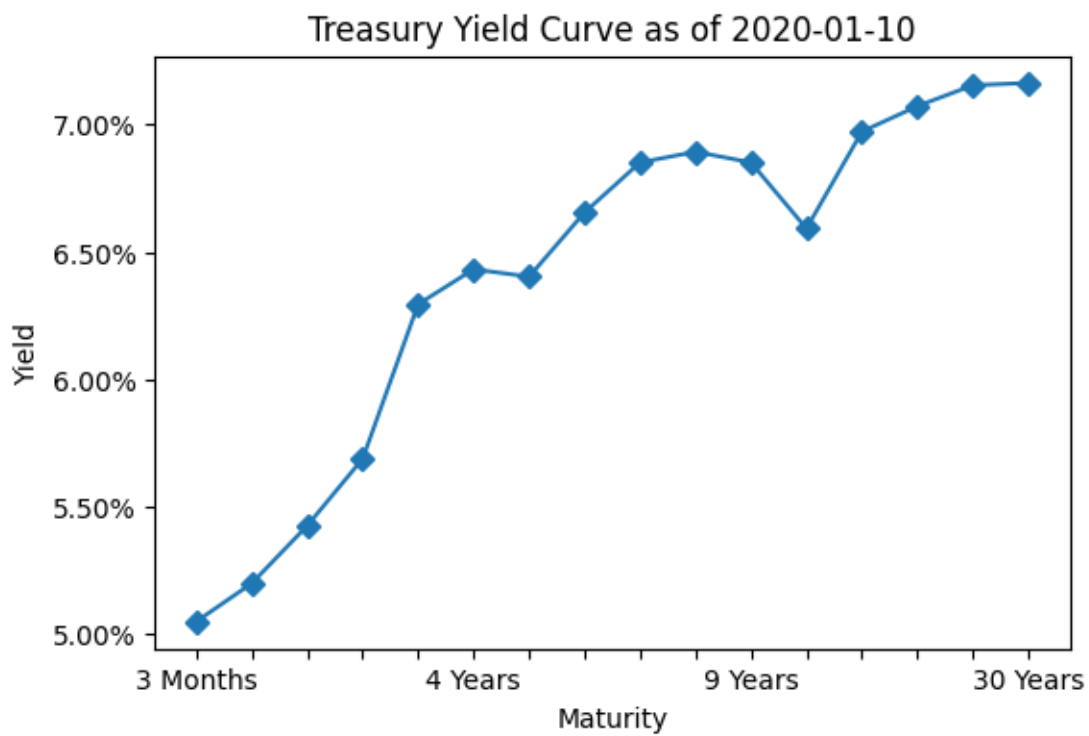
- `t` = maturities
- `y` = yields as of a specific date
- `bc_type=((2, 0.0), (2, 0.0))` sets the **natural end condition** by specifying that the second derivative at the two endpoints is zero.

```
[12]: def fit_cs(t, y, t_hat):
      """
      Fitting a Cubic Spline Model on all the maturities on a specific date.
      """
      cs = CubicSpline(t, y, bc_type=((2, 0.0), (2, 0.0)))
      interpolated_yields = cs(t_hat)
      plt.plot(t_hat, interpolated_yields)
```

```
[13]: date = "2020-01-10"
      t = np.array([0.25,0.5,1,2,3,4,5,6,7,8,9,10,12,15,24,30])
      y = np.array(df.loc[date])
      t_hat = np.linspace(0.5,30,100)
      plot_yield_curve(date, '')
      fit_cs(t, y, t_hat)
```

<ipython-input-8-f6116200d263>:10: UserWarning: `set_ticklabels()` should only be used with a fixed number of ticks, i.e. after `set_ticks()` or using a `FixedLocator`.

```
ax.set_yticklabels(['{:.2f}%'.format(y) for y in ax.get_yticks()])
```



### 1.0.5 Q2. e.

Fit: - NS Model is a parametric model with defined parameters  $\beta_0, \beta_1$  and  $\beta_2$ . The model smoothened the trend of the yield curve and just fit the general level with the slope and curvature of the curve. It failed to capture the nuances of the mid term maturities. Features like a lower yield in the 10 year maturity did not get captured in the NS Model.

- Cubic Spline is a non parametric model which fits a spline between each point individually and thus captures a much more nuanced view of the yield curve. The fit is indeed much better than the NS Model.

Interpretation: - Since NS is a parametric approach, its parameters have a strict economic interpretation. -  $\beta_0$  : Level -  $\beta_1$  : Slope -  $\beta_2$  : Curvature -  $\tau$  : Decay Rate

It is helpful in making economic decision based on these parameters.

- Cubic Spline takes a non parametric approach and as a result does not have a direct economic interpretation. Its use cases are when one needs a data driven approach to mapping the yield movements, with a certain degree of smoothing involved. However, It is very susceptible to overfitting. Thus one must be careful towards the model complexity-overfitting tradeoff.

### 1.0.6 Q2. g.

Indeed Smoothing data can be considered unethical, as discussed in M2 L4. However, whether or not the NS model smoothing is unethical, depends on the specific use cases.

NS Model is a parametric approach to smoothen out the yield curve based on its broad economic parameters,  $\beta_0, \beta_1, \beta_2$  and  $\tau$ . It aims to provide a simplified picture of the economic conditions based on the yields, to provide insights on the expected trends in the market both in the short and long term. In this case the smoothing out is not done to hide variability, it is done to cancel out the noise to identify trends. It is not unethical to smoothen the curve via NS Model in this case.

However, say in a situation where the user does the simplification, but fails to disclose it, with a purpose of hiding variability, or hiding a sharp downturn in the yields, it becomes unethical.

## 2 Task 3

Q3. a & b) Generating 5 Gaussian uncorrelated random variables and running Principal Components on Covariance Matrix

```
[14]: np.random.seed(42)
sim_yld = np.random.normal(loc = 0, scale = 0.01, size = (200,5))
yld = pd.DataFrame(sim_yld, columns = ['3 Months', '1 Year', '3 Years', '10_
↳Years', '30 Years'])
mean_yld = yld.mean()
std_yld = yld.std()
yc_standardised = (yld - mean_yld) / std_yld
std_data_cov = yc_standardised.cov()
eigenvalues, eigenvectors = LA.eig(std_data_cov)
```

```

principal_components = yc_standardised.dot(eigenvectors)
principal_components.columns = ["PC_1", "PC_2", "PC_3", "PC_4", "PC_5"]
df_eigval = pd.DataFrame({"Eigenvalues": eigenvalues}, index=range(1,6))
df_eigval["Explained proportion"] = df_eigval["Eigenvalues"] / np.
    ↪sum(df_eigval["Eigenvalues"])
df_eigval.style.format({"Explained proportion": "{:.2%}"})

```

[14]: <pandas.io.formats.style.Styler at 0x7d09633dc790>

Q3. c) The above df shows the principal components described by their eigen values and the amount of explanation of the variance they respectively do. As we can see that in case of random variables, which are uncorrelated, the 5 principal components more or less describe an equal proportion of the variance. This can be owed to the fact that the multivariate distributions are i.i.d across the observations.

Q3 d) Plotting the Scree plot of the variance explained by each component.

```

[15]: # Plotting the Scree plot

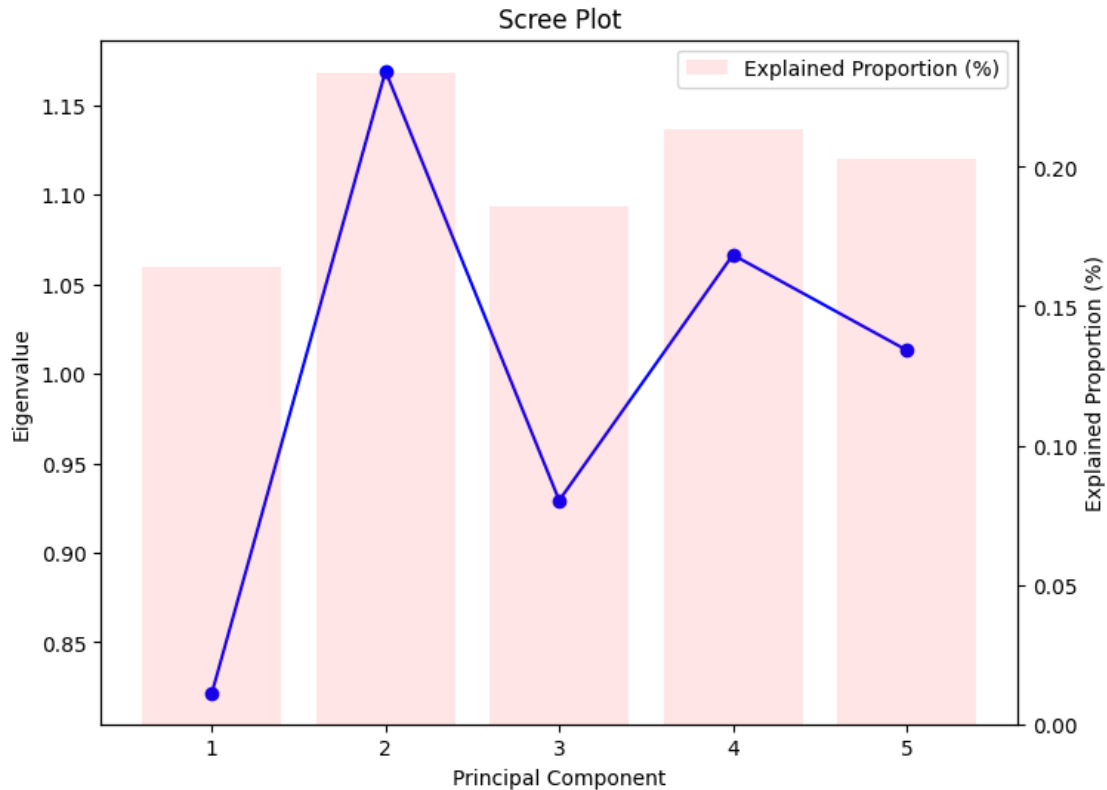
plt.figure(figsize=(8, 6))
plt.plot(range(1, 6), df_eigval['Eigenvalues'], marker='o', linestyle='-',
    ↪color='b', label='Eigenvalues')
plt.title("Scree Plot")
plt.xlabel("Principal Component")
plt.ylabel("Eigenvalue")
plt.xticks(range(1, 6))

# Step 2: Plotting explained proportion as bars

plt.twinx()
plt.bar(range(1, 6), df_eigval['Explained proportion'], alpha=0.1, color='r',
    ↪label="Explained Proportion (%)")
plt.ylabel("Explained Proportion (%)")

plt.grid(False)
plt.legend(loc="upper right")
plt.show()

```



Here we can see that the explained variance does not reduce as the number of principal components increase. Infact they explain more or less similar proportions of the variance.

Q3. e & f) Loading actual data of government securities, for the past 6 months and converting yields into yield changes.

```
[17]: df = pd.read_csv('C:/Users/Gast01/WQU-GWP1/India Bond Yield Data.csv')
df.index = pd.to_datetime(df['Date'])
df_sec = df[df.index >= (df.index.max() - pd.DateOffset(months=6))]

# Selecting 5 securities
df_sec = df_sec[['3 Months', '1 Year', '3 Years', '10 Years', '30 Years']]

# Converting yields into yield changes and dropping NaNs
df_pct_chg = df_sec.pct_change().dropna()
```

Q3. g) Running Principal Components on original yields, using covariance matrix.

```
[18]: yld_chg_mean = df_pct_chg.mean()
yld_chg_std = df_pct_chg.std()
standardized_yld_chag = (df_pct_chg - yld_chg_mean) / yld_chg_std
std_yld_chg_cov = standardized_yld_chag.cov()
eigenvalues_yld, eigenvectors_yld = LA.eig(std_yld_chg_cov)
```

```
principal_components = standardized_yld_chag.dot(eigenvectors_yld)
principal_components.columns = ["PC_1", "PC_2", "PC_3", "PC_4", "PC_5"]
df_eigval_yld = pd.DataFrame({"Eigenvalues": eigenvalues_yld}, index=range(1,6))
df_eigval_yld["Explained proportion"] = df_eigval_yld["Eigenvalues"] / np.
    ↪sum(df_eigval_yld["Eigenvalues"])
df_eigval_yld.style.format({"Explained proportion": "{:.2%}"})
```

[18]: <pandas.io.formats.style.Styler at 0x7d0965c5bf10>

Q3. h) From the above df we can see a different picture regarding the eigenvectors and their explanation of the variance. The first PC explains more than 50% of the variance in the yields. This is the “Level” of the yield curve. This shows the long term expectations of the yields. The second PC explains almost 21% of the variance. We can interpret this as the “Slope” of the yield curve. Finally the 3rd PC explains about 15% of the variance, and this can be attributed to the “Curvature” of the yield curve, or how fast the slope is changing. The last 2 PCs explain relatively lower proportion of the variance, since the level, slope and curvature are the 3 main principal components.

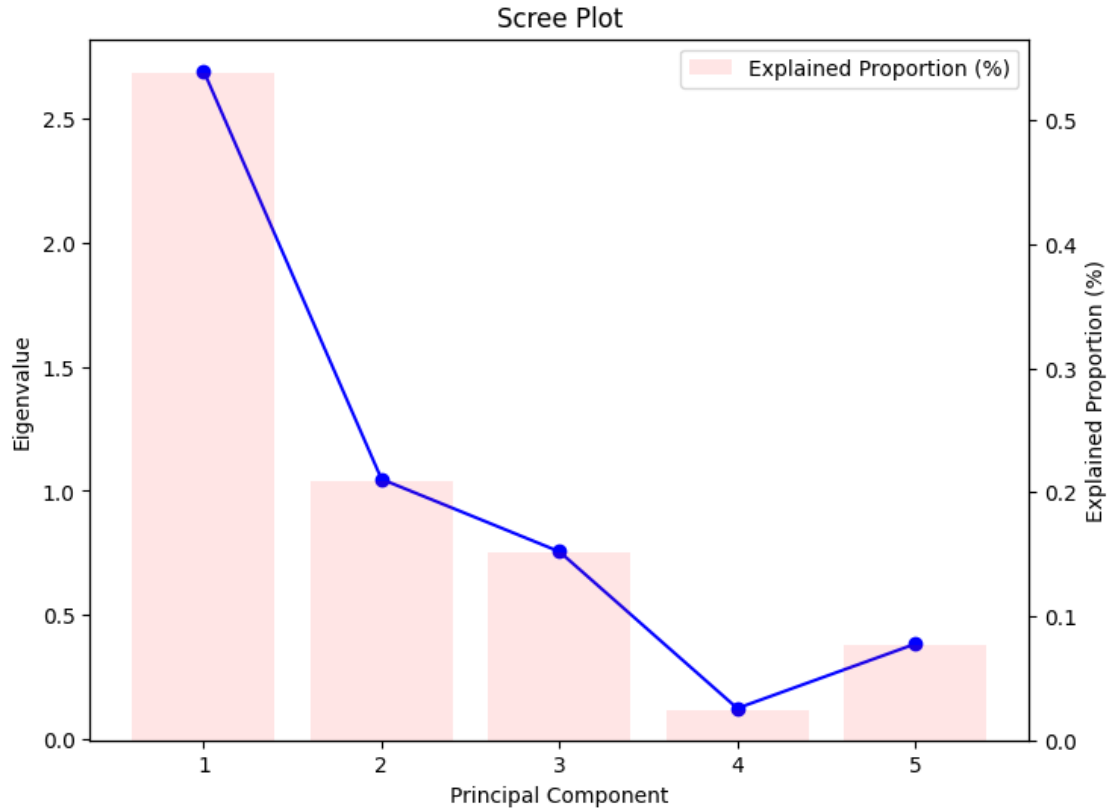
Q3. i) Producing Scree Plot

[19]: *# Plotting the Scree plot*

```
plt.figure(figsize=(8, 6))
plt.plot(range(1, 6), df_eigval_yld['Eigenvalues'], marker='o', linestyle='-',
    ↪color='b', label='Eigenvalues')
plt.title("Scree Plot")
plt.xlabel("Principal Component")
plt.ylabel("Eigenvalue")
plt.xticks(range(1, 6))
plt.grid(False)

# Plotting explained proportion as bars
plt.twinx()
plt.bar(range(1, 6), df_eigval_yld['Explained proportion'], alpha=0.1,
    ↪color='r', label="Explained Proportion (%)")
plt.ylabel("Explained Proportion (%)")
plt.legend(loc="upper right")

# Show the plot
plt.show()
```



Q3. j.

The above scree plot is more in line with our expectations. The plot can be read as an elbow method. The elbow forming when the first principal component is explaining the highest variance and as the number of PCs increase, they explain lesser and lesser of the variance of the yield curve. In our case, the first three principal components explain about 92% of the data.

This observation is different than the uncorrelated gaussian random variables case because the data was generated such that the random variables had a same variance. As a result there is no dominant direction, and the variance is equally spread across the 5 eigenvalues, so the PC algorithm finds similar eigenvalues for the 5 PCs.

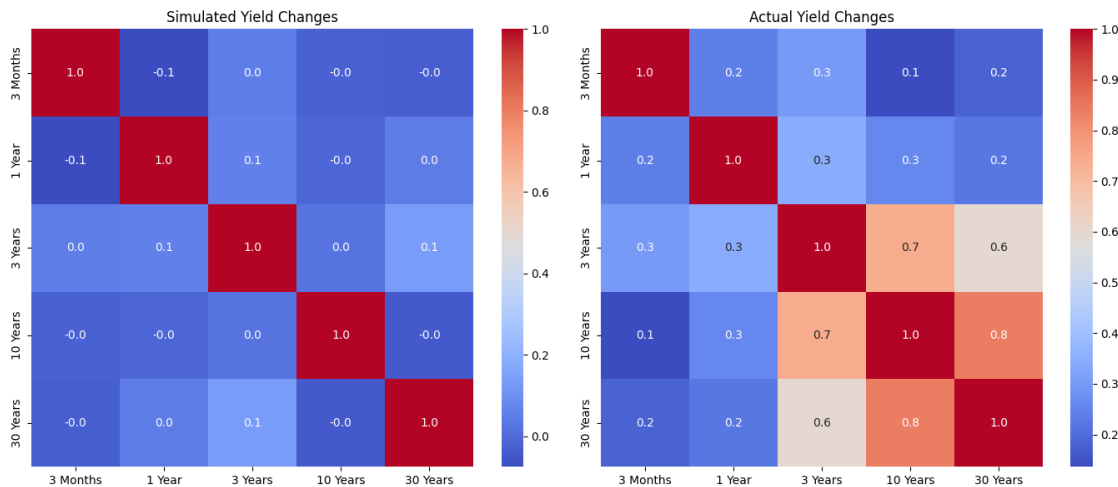
Since principal components are dimensionality reduction techniques, they work well when the variables are closely correlated. From the following figure we see that the yield changes from the actual data is more correlated than the gaussian data generation process. Hence PCA performed better on the real data.

```
[20]: fig, axes = plt.subplots(1, 2, figsize=(14, 6))

# Simulated yield changes
sns.heatmap(yc_standardised.corr(), annot=True, cmap='coolwarm', fmt=".1f",
            ↪ax=axes[0])
axes[0].set_title('Simulated Yield Changes')
```

```
# Actual yield changes
sns.heatmap(standardized_yld_chag.corr(), annot=True, cmap='coolwarm', fmt="
↪1f", ax=axes[1])
axes[1].set_title('Actual Yield Changes')

plt.tight_layout()
plt.show()
```



## QUESTION 4

### Empirical Analysis of ETFs

#### QUESTIONS 4(a):

#### ANSWER

For the group work, we decided to select a utility ETF called XLU, which has about 30 Holdings. The index includes securities of companies from the following industries: electric utilities; water utilities; multi-utilities; independent power and renewable electricity producers; and gas utilities. The fund is non-diversified.

```
[21]: # Importing required libraries

import datetime
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import seaborn as sns
import yfinance as yfin
import math
```



```

from numpy import linalg as LA

from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.decomposition import TruncatedSVD

from datetime import date

pd.options.display.float_format = "{:,.6f}".format

```

**QUESTION 4(b):** Get at least 6 months of data (~ 120 data points).

One year of xlu data had been fetched from yfinance database as follows:

```

[22]: # Starting and end dates
start = datetime.date(2019, 1, 1)
end = datetime.date(2020, 1, 1)

xlu_tickers = [
    "NEE", # NextEra Energy Inc.
    "SO", # Southern Company
    "DUK", # Duke Energy Corporation
    "AEP", # American Electric Power Company Inc.
    "D", # Dominion Energy Inc.
    "SRE", # Sempra
    "EXC", # Exelon Corporation
    "VST", # Vistra Corp.
    "PEG", # Public Service Enterprise Group Inc.
    "PCG", # PG&E Corporation
    "XEL", # Xcel Energy Inc.
    "ED", # Consolidated Edison Inc.
    "AWK", # American Water Works Company Inc.
    "WEC", # WEC Energy Group Inc.
    "EIX", # Edison International
    "ES", # Eversource Energy
    "ATO", # Atmos Energy Corporation
    "CMS", # CMS Energy Corporation
    "NI", # NiSource Inc.
    "PNW", # Pinnacle West Capital Corporation
    "CNP", # CenterPoint Energy Inc.
    "EVRG", # Evergy Inc.
    "FE", # FirstEnergy Corp.
    "NRG", # NRG Energy Inc.
    "OGE", # OGE Energy Corp.
    "AEE", # Ameren Corporation
    "AES", # The AES Corporation
    "LNT", # Alliant Energy Corporation
    "UGI", # UGI Corporation

```

```

    "IDA"      # IDACORP Inc.
]

# Get ETF data
df = yfin.download(xlu_tickers, start, end, auto_adjust = False)["Adj Close"]

# Convert DataFrame index to timezone-aware (UTC)
df.index = df.index.tz_localize('UTC')

```

[\*\*\*\*\*100%\*\*\*\*\*] 30 of 30 completed

Let us have a look at the first five rows of the daily data.

[23]: df.head(5)

```

[23]: Ticker          AEE          AEP          AES          ATO          AWK  \
Date
2019-01-02 00:00:00+00:00 53.270641 58.418915 11.534038 77.131935 79.253159
2019-01-03 00:00:00+00:00 53.404797 58.282764 11.525901 77.715683 79.565666
2019-01-04 00:00:00+00:00 54.159466 58.819359 11.908203 78.711479 80.163826
2019-01-07 00:00:00+00:00 53.538967 58.490982 11.965139 78.136322 79.315689
2019-01-08 00:00:00+00:00 54.385860 59.211826 12.241695 79.157890 80.476295

```

```

Ticker          CMS          CNP          D          DUK          ED  \
Date
2019-01-02 00:00:00+00:00 40.172016 23.255310 54.124611 65.633408 59.839409
2019-01-03 00:00:00+00:00 40.246845 23.388199 53.866074 65.610184 59.990723
2019-01-04 00:00:00+00:00 40.621037 23.820086 54.907814 66.152603 60.962288
2019-01-07 00:00:00+00:00 40.413155 23.911442 54.375546 65.873627 60.078320
2019-01-08 00:00:00+00:00 41.020195 24.368248 55.006657 66.702774 60.444645

```

```

Ticker          ...          OGE          PCG          PEG          PNW  \
Date          ...
2019-01-02 00:00:00+00:00  ... 28.728683 23.682159 41.137543 63.742741
2019-01-03 00:00:00+00:00  ... 28.983988 23.831419 41.129467 64.388840
2019-01-04 00:00:00+00:00  ... 29.442028 24.279190 41.574581 65.181107
2019-01-07 00:00:00+00:00  ... 29.284349 18.856173 41.663612 65.058037
2019-01-08 00:00:00+00:00  ... 29.967651 17.473055 41.825474 66.034920

```

```

Ticker          SO          SRE          UGI          VST          WEC  \
Date
2019-01-02 00:00:00+00:00 33.987125 43.674419 40.206642 19.215071 55.447842
2019-01-03 00:00:00+00:00 34.469093 43.874878 40.706715 19.034771 55.595531
2019-01-04 00:00:00+00:00 34.756718 45.188049 41.591484 19.807493 56.005760
2019-01-07 00:00:00+00:00 34.678986 45.179878 41.676121 20.142340 55.751411
2019-01-08 00:00:00+00:00 35.666275 46.443966 41.991550 20.150928 56.325771

```

```

Ticker          XEL

```

```
Date
2019-01-02 00:00:00+00:00 39.967186
2019-01-03 00:00:00+00:00 39.809719
2019-01-04 00:00:00+00:00 40.199272
2019-01-07 00:00:00+00:00 40.025204
2019-01-08 00:00:00+00:00 40.489357
```

[5 rows x 30 columns]

## Overview of the ETF Data

We can use the pandas `describe()` method to show summary stats of our data. We can see that all assets have same number of observations (count) since they all belong to the same portfolio. The other summary stats are relatively basic, like mean and standard deviation along with showing minimum, maximum, and a few quantiles.

```
[24]: df.describe()
```

```
[24]: Ticker      AEE      AEP      AES      ATO      AWK      CMS  \
count  252.000000  252.000000  252.000000  252.000000  252.000000  252.000000
mean    62.849828   71.180582   14.031441   90.896079  101.687287   49.038936
std      3.245782    5.109215    1.068596    5.416583    9.813963    3.909191
min     53.270641   58.282764   11.525901   77.131935   79.253159   40.172016
25%     60.837073   67.466379   13.314819   87.135109   93.918928   45.967123
50%     63.895060   72.929676   13.999063   92.650085  104.843285   49.257168
75%     65.094200   75.268436   14.701836   95.203949  110.105621   52.592816
max     68.704857   78.527794   16.838137   99.922478  117.211868   55.235107

Ticker      CNP      D      DUK      ED  ...      OGE      PCG  \
count  252.000000  252.000000  252.000000  252.000000  ...  252.000000  252.000000
mean    24.547178   59.984373   71.137369   69.934112  ...   32.676920   14.893234
std      1.196644    3.548472    2.900220    4.330082  ...    1.146570    5.269217
min     21.099731   51.683754   64.827545   59.289909  ...   28.728683    3.781185
25%     24.072486   58.064060   69.467724   67.827486  ...   32.078154   10.634583
50%     24.733864   59.510866   70.553871   70.953808  ...   32.867317   17.030258
75%     25.513127   63.671901   72.968153   72.358799  ...   33.288613   18.866123
max     26.322069   66.020332   77.733910   77.404655  ...   35.019447   24.279190

Ticker      PEG      PNW      SO      SRE      UGI      VST  \
count  252.000000  252.000000  252.000000  252.000000  252.000000  252.000000
mean    48.428759   72.444122   44.384742   55.560343   39.517781   21.618953
std      2.487902    3.038539    4.653334    5.162168    2.706561    1.422285
min     41.072807   63.435051   33.987125   43.674419   32.437710   18.584152
25%     47.887318   70.596058   40.785181   51.892804   37.717446   20.698181
50%     48.764412   73.282558   44.307463   56.768837   40.448433   21.774376
75%     50.065631   74.517906   49.139084   59.809653   41.575361   22.810292
max     52.458248   77.214256   51.908272   63.737045   43.899559   23.848814
```

| Ticker | WEC        | XEL        |
|--------|------------|------------|
| count  | 252.000000 | 252.000000 |
| mean   | 70.069637  | 49.470813  |
| std    | 6.990331   | 3.973183   |
| min    | 55.447842  | 39.809719  |
| 25%    | 64.181870  | 46.539076  |
| 50%    | 71.049873  | 50.785336  |
| 75%    | 76.504229  | 52.605731  |
| max    | 81.964592  | 55.469288  |

[8 rows x 30 columns]

#### QUESTION 4(C) :

#### ANSWER Importance of Daily Returns

In financial analysis, the daily return of an asset, XLU-ETF in our case, is a fundamental metric used to measure its performance over a single trading day. It definitely aids in quantifying the percentage change in the asset's value between the close of one trading day and the close of the subsequent day. To compute a daily return you have to use the closing prices of the asset on two consecutive trading days.

To account for dividend and stock splits, technically, one must use adjusted closing price to avoid impact on the individual prices of an ETF constituents assets/holdings.

```
[25]: # 1. Compute daily log returns
daily_returns = df.pct_change() # Used simple return to compute the daily
    ↪return (alternatively: log return can also be used, np.log(df) - np.log(df.
    ↪shift(1))
daily_returns = daily_returns.dropna()
daily_returns.head()
```

```
[25]: Ticker          AEE          AEP          AES          ATO          AWK  \
Date
2019-01-03 00:00:00+00:00  0.002518 -0.002331 -0.000705  0.007568  0.003943
2019-01-04 00:00:00+00:00  0.014131  0.009207  0.033169  0.012813  0.007518
2019-01-07 00:00:00+00:00 -0.011457 -0.005583  0.004781 -0.007307 -0.010580
2019-01-08 00:00:00+00:00  0.015818  0.012324  0.023113  0.013074  0.014633
2019-01-09 00:00:00+00:00 -0.007401 -0.007575 -0.001993 -0.016376 -0.012758
```

```
Ticker          CMS          CNP          D          DUK          ED  \
Date
2019-01-03 00:00:00+00:00  0.001863  0.005714 -0.004777 -0.000354  0.002529
2019-01-04 00:00:00+00:00  0.009297  0.018466  0.019339  0.008267  0.016195
2019-01-07 00:00:00+00:00 -0.005118  0.003835 -0.009694 -0.004217 -0.014500
2019-01-08 00:00:00+00:00  0.015021  0.019104  0.011607  0.012587  0.006097
2019-01-09 00:00:00+00:00 -0.009528 -0.012270 -0.001659 -0.014754 -0.007905
```

```
Ticker          ...          OGE          PCG          PEG          PNW  \
```

| Date                      | ... |           |           |           |           |  |
|---------------------------|-----|-----------|-----------|-----------|-----------|--|
| 2019-01-03 00:00:00+00:00 | ... | 0.008887  | 0.006303  | -0.000196 | 0.010136  |  |
| 2019-01-04 00:00:00+00:00 | ... | 0.015803  | 0.018789  | 0.010822  | 0.012304  |  |
| 2019-01-07 00:00:00+00:00 | ... | -0.005356 | -0.223361 | 0.002141  | -0.001888 |  |
| 2019-01-08 00:00:00+00:00 | ... | 0.023333  | -0.073351 | 0.003885  | 0.015016  |  |
| 2019-01-09 00:00:00+00:00 | ... | -0.001897 | 0.015376  | -0.007546 | -0.013861 |  |

| Ticker                    |  | SO        | SRE       | UGI       | VST       | WEC       | \ |
|---------------------------|--|-----------|-----------|-----------|-----------|-----------|---|
| Date                      |  |           |           |           |           |           |   |
| 2019-01-03 00:00:00+00:00 |  | 0.014181  | 0.004590  | 0.012438  | -0.009383 | 0.002664  |   |
| 2019-01-04 00:00:00+00:00 |  | 0.008344  | 0.029930  | 0.021735  | 0.040595  | 0.007379  |   |
| 2019-01-07 00:00:00+00:00 |  | -0.002236 | -0.000181 | 0.002035  | 0.016905  | -0.004541 |   |
| 2019-01-08 00:00:00+00:00 |  | 0.028469  | 0.027979  | 0.007569  | 0.000426  | 0.010302  |   |
| 2019-01-09 00:00:00+00:00 |  | -0.008501 | -0.009777 | -0.010260 | -0.006817 | -0.005973 |   |

| Ticker                    | XEL       |
|---------------------------|-----------|
| Date                      |           |
| 2019-01-03 00:00:00+00:00 | -0.003940 |
| 2019-01-04 00:00:00+00:00 | 0.009785  |
| 2019-01-07 00:00:00+00:00 | -0.004330 |
| 2019-01-08 00:00:00+00:00 | 0.011597  |
| 2019-01-09 00:00:00+00:00 | -0.007983 |

[5 rows x 30 columns]

#### Question 4(d): Compute the covariance matrix.

Before standardizing the data, you need to calculate the daily returns of the stocks in the XLU ETF, which is typically done using the percentage change between consecutive closing prices. The covariance matrix is calculated using the daily returns, not the raw prices. It measures how the returns of different XLU ETF stocks move together.

Without calculating the daily returns first, the standardization wouldn't make sense. In fact, standardization makes varying swings (big and small) comparable by adjusting them so they all have an average change of 0 and typical change of 1. This is helpful for analysis where the scale of the data matters.

```
[26]: # Standardize stock returns dataset
daily_returns_means = daily_returns.mean()
daily_returns_stds = daily_returns.std()
standardized_returns = (daily_returns - daily_returns_means) / \
    ↪daily_returns_stds
standardized_returns.head()
```

| Ticker                    |  | AEE       | AEP       | AES       | ATO       | AWK       | \ |
|---------------------------|--|-----------|-----------|-----------|-----------|-----------|---|
| Date                      |  |           |           |           |           |           |   |
| 2019-01-03 00:00:00+00:00 |  | 0.183450  | -0.417105 | -0.180855 | 0.763051  | 0.275034  |   |
| 2019-01-04 00:00:00+00:00 |  | 1.498006  | 0.948051  | 2.525074  | 1.371657  | 0.663016  |   |
| 2019-01-07 00:00:00+00:00 |  | -1.398540 | -0.801923 | 0.257431  | -0.962968 | -1.301269 |   |

|                           |           |           |           |           |           |
|---------------------------|-----------|-----------|-----------|-----------|-----------|
| 2019-01-08 00:00:00+00:00 | 1.688990  | 1.316904  | 1.721830  | 1.401925  | 1.435249  |
| 2019-01-09 00:00:00+00:00 | -0.939365 | -1.037670 | -0.283738 | -2.015199 | -1.537639 |

| Ticker                    |  | CMS       | CNP       | D         | DUK       | ED \      |
|---------------------------|--|-----------|-----------|-----------|-----------|-----------|
| Date                      |  |           |           |           |           |           |
| 2019-01-03 00:00:00+00:00 |  | 0.074791  | 0.474443  | -0.621066 | -0.107618 | 0.196326  |
| 2019-01-04 00:00:00+00:00 |  | 0.908490  | 1.557158  | 2.049130  | 0.984212  | 1.857444  |
| 2019-01-07 00:00:00+00:00 |  | -0.707954 | 0.314894  | -1.165504 | -0.596891 | -1.873469 |
| 2019-01-08 00:00:00+00:00 |  | 1.550297  | 1.611336  | 1.192923  | 1.531274  | 0.630099  |
| 2019-01-09 00:00:00+00:00 |  | -1.202506 | -1.052580 | -0.275826 | -1.931287 | -1.071834 |

| Ticker                    | ... | OGE       | PCG       | PEG       | PNW \     |
|---------------------------|-----|-----------|-----------|-----------|-----------|
| Date                      | ... |           |           |           |           |
| 2019-01-03 00:00:00+00:00 | ... | 0.936070  | 0.058488  | -0.115138 | 1.020381  |
| 2019-01-04 00:00:00+00:00 | ... | 1.734418  | 0.196466  | 1.214921  | 1.250011  |
| 2019-01-07 00:00:00+00:00 | ... | -0.707878 | -2.479326 | 0.167058  | -0.252974 |
| 2019-01-08 00:00:00+00:00 | ... | 2.603599  | -0.821695 | 0.377514  | 1.537120  |
| 2019-01-09 00:00:00+00:00 | ... | -0.308636 | 0.158749  | -1.002372 | -1.520917 |

| Ticker                    |  | SO        | SRE       | UGI       | VST       | WEC \     |
|---------------------------|--|-----------|-----------|-----------|-----------|-----------|
| Date                      |  |           |           |           |           |           |
| 2019-01-03 00:00:00+00:00 |  | 1.412399  | 0.330930  | 0.903817  | -0.718701 | 0.143397  |
| 2019-01-04 00:00:00+00:00 |  | 0.750993  | 3.088814  | 1.559423  | 2.999632  | 0.675746  |
| 2019-01-07 00:00:00+00:00 |  | -0.448059 | -0.188285 | 0.170297  | 1.237110  | -0.670053 |
| 2019-01-08 00:00:00+00:00 |  | 3.031608  | 2.876486  | 0.560491  | 0.011123  | 1.005792  |
| 2019-01-09 00:00:00+00:00 |  | -1.157927 | -1.232719 | -0.696680 | -0.527805 | -0.831618 |

| Ticker                    |  | XEL       |
|---------------------------|--|-----------|
| Date                      |  |           |
| 2019-01-03 00:00:00+00:00 |  | -0.582390 |
| 2019-01-04 00:00:00+00:00 |  | 0.959800  |
| 2019-01-07 00:00:00+00:00 |  | -0.626237 |
| 2019-01-08 00:00:00+00:00 |  | 1.163302  |
| 2019-01-09 00:00:00+00:00 |  | -1.036720 |

[5 rows x 30 columns]

```
[27]: # Calculate covariance for standardized return matrix
standardized_returns_dvd_sqrt_n=(standardized_returns/math.
    ↳sqrt(len(standardized_returns)-1))
standardized_returns_cov = standardized_returns_dvd_sqrt_n.
    ↳T@standardized_returns_dvd_sqrt_n
standardized_returns_cov.head()
```

| [27]: Ticker | AEE      | AEP      | AES      | ATO      | AWK      | CMS      | CNP \    |
|--------------|----------|----------|----------|----------|----------|----------|----------|
| Ticker       |          |          |          |          |          |          |          |
| AEE          | 1.000000 | 0.806508 | 0.522383 | 0.771703 | 0.724853 | 0.834286 | 0.400582 |

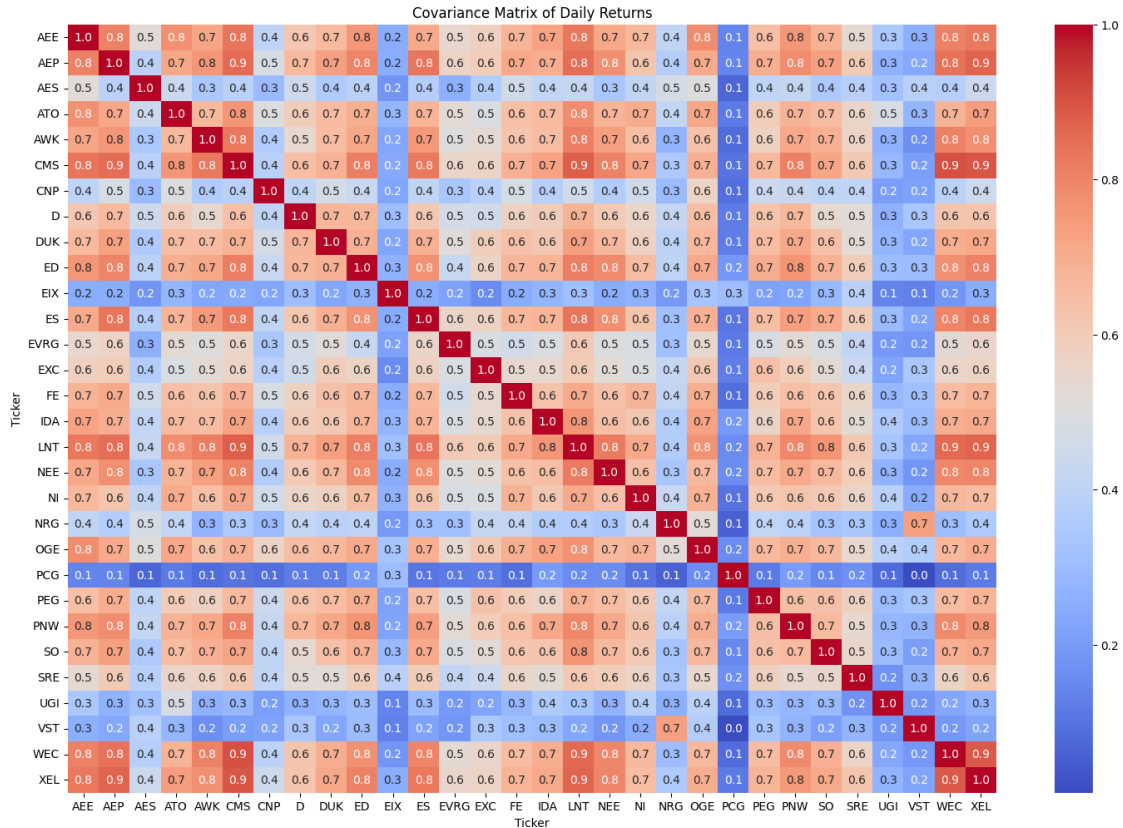
|     |          |          |          |          |          |          |          |
|-----|----------|----------|----------|----------|----------|----------|----------|
| AEP | 0.806508 | 1.000000 | 0.389617 | 0.716671 | 0.757821 | 0.851061 | 0.479604 |
| AES | 0.522383 | 0.389617 | 1.000000 | 0.440487 | 0.343414 | 0.386723 | 0.291766 |
| ATO | 0.771703 | 0.716671 | 0.440487 | 1.000000 | 0.665880 | 0.763396 | 0.497985 |
| AWK | 0.724853 | 0.757821 | 0.343414 | 0.665880 | 1.000000 | 0.827253 | 0.360149 |

| Ticker | D        | DUK      | ED       | ... | OGE      | PCG      | PEG      | PNW      | \ |
|--------|----------|----------|----------|-----|----------|----------|----------|----------|---|
| Ticker |          |          |          | ... |          |          |          |          |   |
| AEE    | 0.620391 | 0.662887 | 0.757679 | ... | 0.781366 | 0.138427 | 0.646135 | 0.758153 |   |
| AEP    | 0.661853 | 0.735293 | 0.800920 | ... | 0.712110 | 0.143775 | 0.699680 | 0.782281 |   |
| AES    | 0.451335 | 0.354717 | 0.390818 | ... | 0.505822 | 0.070451 | 0.406285 | 0.421399 |   |
| ATO    | 0.592841 | 0.651968 | 0.697263 | ... | 0.720000 | 0.113871 | 0.600165 | 0.683688 |   |
| AWK    | 0.540264 | 0.659173 | 0.713845 | ... | 0.638203 | 0.082087 | 0.575091 | 0.716895 |   |

| Ticker | SO       | SRE      | UGI      | VST      | WEC      | XEL      |
|--------|----------|----------|----------|----------|----------|----------|
| Ticker |          |          |          |          |          |          |
| AEE    | 0.689448 | 0.538975 | 0.345586 | 0.258085 | 0.808670 | 0.827982 |
| AEP    | 0.734963 | 0.604644 | 0.263486 | 0.235226 | 0.840926 | 0.866046 |
| AES    | 0.350035 | 0.378619 | 0.316272 | 0.381778 | 0.376726 | 0.441356 |
| ATO    | 0.652024 | 0.559601 | 0.485826 | 0.279782 | 0.686117 | 0.742385 |
| AWK    | 0.691632 | 0.566794 | 0.261875 | 0.164899 | 0.810270 | 0.779339 |

[5 rows x 30 columns]

```
[28]: plt.figure(figsize=(18, 12))
sns.heatmap(standardized_returns_cov, annot=True, cmap='coolwarm', fmt=".1f")
plt.title('Covariance Matrix of Daily Returns')
plt.show()
```



#### Question 4(e):

(compare and contrast PCA and SVD, explain what the eigenvectors, eigenvalues, singular values etc show us for the specific data, etc,)

#### ANSWER

Both PCA and SVD are dimensionality reduction technique. They reduce dimensions of matrices but importantly they do retain substantial information or rather variance of a given financial data.

PCA plays a critical role of identifying key factors behind asset price movements whilst reducing dimensionality of risk-based models. Through PCA, one can establish a diversified portfolio based on the principal components. PCA can keep track of unique patterns and ambiguities in portfolio/ETF, which might not be captured in individual asset.

SVD stands for Singular Value Decomposition. It is a robust factorization technique that can effectively decompose any given matrix (Symmetric or not) into 3 matrices, namely, U, S and  $V^T$ . These matrices have  $m \times n$ ,  $m \times n$  and  $n \times n$  dimensions respectively. The magnitude of each singular value shows the importance of the corresponding dimension in the data.

Eigen vectors are extracted from PCA and they do represent the principal components which depict relative movements in daily returns of assets. On the other hand, eigenvalues, provides the amount of variance in the daily returns data related to principal components (eigenvector). The higher the eigenvalue, the better it gives information about the asset under analysis.



The following code solutions, shows how Eigenvalues and Eigenvectors help generate Principal Components based on ULX-ETF data under analysis. It is important to note that SVD can be applied to generate eigenvalues and eigenvectors.

```
[29]: # Calculate eigenvectors and eigenvalues of the covariance matrix of
      ↪standardized dataset
      eigenvalues, eigenvectors = np.linalg.eig(standardized_returns_cov)
      eigenvalues
```

```
[29]: array([16.94028088,  1.84467656,  1.26660112,  1.0033836 ,  0.85284607,
            0.78526804,  0.68313883,  0.62148367,  0.56747695,  0.54817649,
            0.50067321,  0.4355354 ,  0.39437057,  0.06999438,  0.08276568,
            0.09744171,  0.10123079,  0.13109828,  0.36516273,  0.35232383,
            0.33001439,  0.29732338,  0.28207821,  0.16274534,  0.17443758,
            0.19353803,  0.24393001,  0.23422616,  0.21796549,  0.21981263])
```

```
[30]: print(pd.DataFrame(eigenvectors).head())
```

|   | 0        | 1         | 2         | 3         | 4         | 5         | 6         | \ |
|---|----------|-----------|-----------|-----------|-----------|-----------|-----------|---|
| 0 | 0.213975 | 0.029617  | 0.048260  | -0.097896 | 0.130891  | 0.023761  | 0.042475  |   |
| 1 | 0.219064 | 0.116065  | 0.034320  | 0.081349  | 0.006026  | -0.016249 | 0.048878  |   |
| 2 | 0.127139 | -0.328994 | 0.058621  | -0.160035 | 0.144884  | 0.216879  | 0.387771  |   |
| 3 | 0.202651 | -0.031364 | -0.007038 | -0.274620 | -0.072991 | 0.025539  | -0.076439 |   |
| 4 | 0.201148 | 0.183299  | 0.061812  | -0.008566 | 0.061080  | 0.105053  | -0.072025 |   |

|   | 7         | 8         | 9         | ... | 20        | 21        | 22        | 23        | \ |
|---|-----------|-----------|-----------|-----|-----------|-----------|-----------|-----------|---|
| 0 | -0.207650 | 0.096307  | 0.051588  | ... | -0.028883 | 0.341933  | -0.220457 | -0.238741 |   |
| 1 | 0.026485  | -0.012477 | 0.066108  | ... | 0.192387  | 0.086952  | -0.160607 | 0.223562  |   |
| 2 | -0.658986 | -0.141501 | -0.186631 | ... | -0.012539 | -0.139527 | 0.114488  | 0.021493  |   |
| 3 | 0.032491  | 0.134994  | 0.117323  | ... | -0.166482 | 0.374353  | -0.255861 | 0.304626  |   |
| 4 | 0.013562  | 0.152696  | -0.084538 | ... | -0.322136 | -0.388158 | 0.055149  | 0.373392  |   |

|   | 24        | 25        | 26        | 27        | 28        | 29        |
|---|-----------|-----------|-----------|-----------|-----------|-----------|
| 0 | -0.037419 | 0.028039  | -0.020517 | -0.062485 | 0.291977  | 0.005230  |
| 1 | 0.065085  | -0.054074 | 0.125435  | 0.118862  | 0.127297  | 0.178015  |
| 2 | -0.085401 | -0.007385 | 0.002403  | -0.010940 | -0.185881 | -0.053974 |
| 3 | 0.157672  | 0.122336  | -0.450374 | -0.105967 | -0.379816 | -0.083053 |
| 4 | -0.050213 | 0.051465  | -0.031854 | -0.431146 | 0.331905  | -0.270027 |

[5 rows x 30 columns]

```
[31]: # Transform standardized data with Loadings
      principal_components = standardized_returns_cov.dot(eigenvectors)
      principal_components.columns = ["PC_" + str(i) for i in range(1, 31)]
      principal_components.head()
```

```
[31]:
```

|        | PC_1 | PC_2 | PC_3 | PC_4 | PC_5 | PC_6 | PC_7 | \ |
|--------|------|------|------|------|------|------|------|---|
| Ticker |      |      |      |      |      |      |      |   |

|     |          |           |           |           |           |           |           |
|-----|----------|-----------|-----------|-----------|-----------|-----------|-----------|
| AEE | 3.624796 | 0.054633  | 0.061126  | -0.098228 | 0.111630  | 0.018659  | 0.029016  |
| AEP | 3.710998 | 0.214102  | 0.043469  | 0.081624  | 0.005139  | -0.012760 | 0.033391  |
| AES | 2.153771 | -0.606888 | 0.074249  | -0.160577 | 0.123564  | 0.170308  | 0.264902  |
| ATO | 3.432957 | -0.057857 | -0.008914 | -0.275549 | -0.062250 | 0.020055  | -0.052218 |
| AWK | 3.407511 | 0.338127  | 0.078291  | -0.008595 | 0.052092  | 0.082494  | -0.049203 |

|        |           |           |           |     |           |           |           |   |
|--------|-----------|-----------|-----------|-----|-----------|-----------|-----------|---|
|        | PC_8      | PC_9      | PC_10     | ... | PC_21     | PC_22     | PC_23     | \ |
| Ticker |           |           |           | ... |           |           |           |   |
| AEE    | -0.129051 | 0.054652  | 0.028280  | ... | -0.009532 | 0.101665  | -0.062186 |   |
| AEP    | 0.016460  | -0.007080 | 0.036239  | ... | 0.063491  | 0.025853  | -0.045304 |   |
| AES    | -0.409549 | -0.080298 | -0.102307 | ... | -0.004138 | -0.041485 | 0.032295  |   |
| ATO    | 0.020193  | 0.076606  | 0.064314  | ... | -0.054941 | 0.111304  | -0.072173 |   |
| AWK    | 0.008429  | 0.086651  | -0.046342 | ... | -0.106310 | -0.115409 | 0.015556  |   |

|        |           |           |           |           |           |           |           |
|--------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|
|        | PC_24     | PC_25     | PC_26     | PC_27     | PC_28     | PC_29     | PC_30     |
| Ticker |           |           |           |           |           |           |           |
| AEE    | -0.038854 | -0.006527 | 0.005427  | -0.005005 | -0.014636 | 0.063641  | 0.001150  |
| AEP    | 0.036384  | 0.011353  | -0.010465 | 0.030597  | 0.027841  | 0.027746  | 0.039130  |
| AES    | 0.003498  | -0.014897 | -0.001429 | 0.000586  | -0.002562 | -0.040516 | -0.011864 |
| ATO    | 0.049576  | 0.027504  | 0.023677  | -0.109860 | -0.024820 | -0.082787 | -0.018256 |
| AWK    | 0.060768  | -0.008759 | 0.009960  | -0.007770 | -0.100986 | 0.072344  | -0.059355 |

[5 rows x 30 columns]

Understanding the explained proportion helps us determine the relative importance of each principal component. By selecting the top principal components with the highest explained proportions, we can reduce the dimensionality of the data while retaining most of the important information.

We used the following codes to demonstrate how to organize the calculated eigenvalues and their corresponding explained proportions into a pandas DataFrame for easier analysis.

Explained Proportion of 56.47% means that this 'PC\_1' alone captures or explains 56.47% of the total variance observed in the standardized daily returns of the XLU ETF's assets.

```
[32]: # Put data into a DataFrame
df_eigval = pd.DataFrame({"Eigenvalues":eigenvalues}, index=range(1,31))

# Work out explained proportion
df_eigval["Explained proportion"] = df_eigval["Eigenvalues"] / np.
    ↳sum(df_eigval["Eigenvalues"])
#Format as percentage
df_eigval.style.format({"Explained proportion": "{:.2%}"})
```

```
[32]: <pandas.io.formats.style.Styler at 0x7d095e4cb810>
```

```
[33]: # Visualization for PCA

x = standardized_returns_cov.values # Convert DataFrame to NumPy array
x = StandardScaler().fit_transform(x) # Standardize the data
```

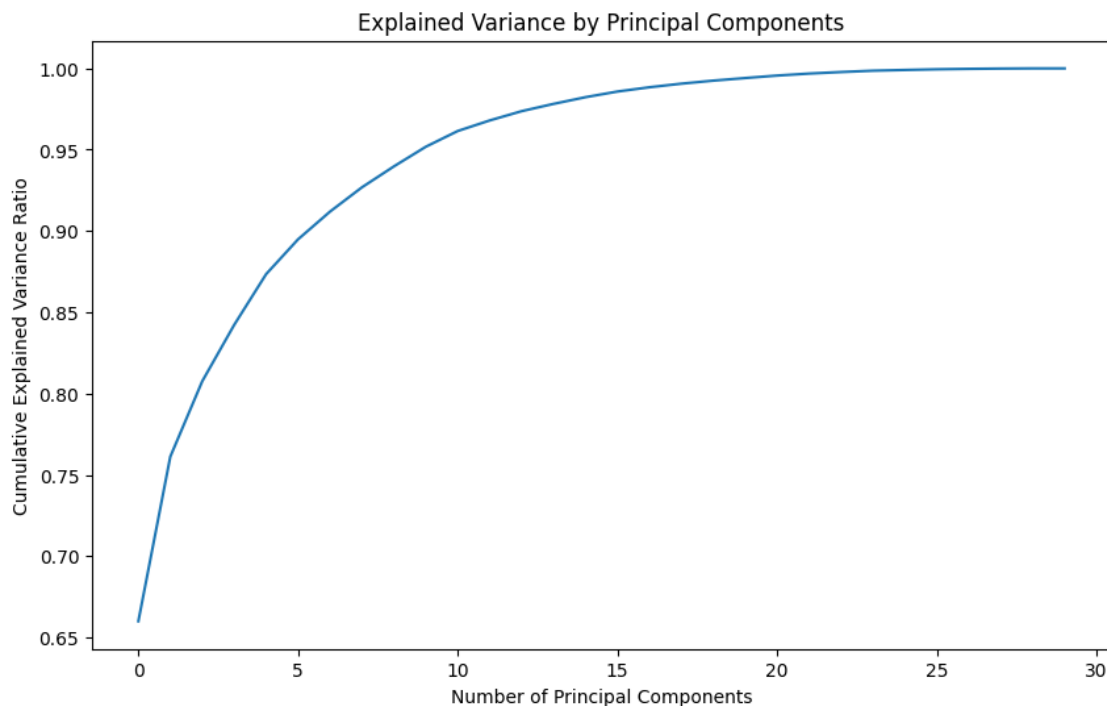
```

# Apply PCA
pca = PCA(n_components=standardized_returns_cov.shape[1])
principalComponents = pca.fit_transform(x)
principalDf = pd.DataFrame(data=principalComponents, columns=['principal_
↪component ' + str(i) for i in range(1, standardized_returns_cov.shape[1] +
↪1)], index=standardized_returns_cov.index)

# Get explained variance ratio
explained_variance_ratio = pca.explained_variance_ratio_

plt.figure(figsize=(10, 6))
plt.plot(np.cumsum(explained_variance_ratio))
plt.xlabel('Number of Principal Components')
plt.ylabel('Cumulative Explained Variance Ratio')
plt.title('Explained Variance by Principal Components')
plt.show()

```



#### Question 4(f): Compute the SVD.

The task here involves decomposing the daily returns matrix of the XLU ETF's assets using Singular Value Decomposition (SVD). SVD is a powerful technique for breaking down a matrix into its fundamental components, revealing underlying relationships and patterns within the data.

By applying SVD to the daily returns matrix, we aim to identify the most significant dimensions

or factors that drive the variability in the returns of the XLU ETF's assets. Secondly, gain insights into the relationships between different assets within the ETF and how they contribute to overall portfolio risk and return and transform the original data into a more manageable form, making it easier to interpret and model.

```
[34]: # Perform SVD for stock returns
      U, s, VT = np.linalg.svd(daily_returns)

[36]: # Use SVD to calculate eigenvectors and eigenvalues of the covariance matrix of
      ↪standardized returns
      U_st_return, s_st_return, VT_st_return = np.linalg.
      ↪svd(standardized_returns_dvd_sqrt_n)
      print("\nSquared Singular values (eigenvalues):")
      print(s_st_return**2)
      print("\nMatrix V (eigenvectors)")
      print(pd.DataFrame(VT_st_return.T).head())
```

Squared Singular values (eigenvalues):

|             |            |            |            |            |             |
|-------------|------------|------------|------------|------------|-------------|
| 16.94028088 | 1.84467656 | 1.26660112 | 1.0033836  | 0.85284607 | 0.78526804  |
| 0.68313883  | 0.62148367 | 0.56747695 | 0.54817649 | 0.50067321 | 0.4355354   |
| 0.39437057  | 0.36516273 | 0.35232383 | 0.33001439 | 0.29732338 | 0.28207821  |
| 0.24393001  | 0.23422616 | 0.21981263 | 0.21796549 | 0.19353803 | 0.17443758  |
| 0.16274534  | 0.13109828 | 0.10123079 | 0.09744171 | 0.08276568 | 0.06999438] |

Matrix V (eigenvectors)

|   | 0         | 1         | 2         | 3         | 4         | 5         | 6         | \ |
|---|-----------|-----------|-----------|-----------|-----------|-----------|-----------|---|
| 0 | -0.213975 | -0.029617 | 0.048260  | -0.097896 | 0.130891  | -0.023761 | 0.042475  |   |
| 1 | -0.219064 | -0.116065 | 0.034320  | 0.081349  | 0.006026  | 0.016249  | 0.048878  |   |
| 2 | -0.127139 | 0.328994  | 0.058621  | -0.160035 | 0.144884  | -0.216879 | 0.387771  |   |
| 3 | -0.202651 | 0.031364  | -0.007038 | -0.274620 | -0.072991 | -0.025539 | -0.076439 |   |
| 4 | -0.201148 | -0.183299 | 0.061812  | -0.008566 | 0.061080  | -0.105053 | -0.072025 |   |

|   | 7         | 8         | 9         | ... | 20        | 21        | 22        | 23        | \ |
|---|-----------|-----------|-----------|-----|-----------|-----------|-----------|-----------|---|
| 0 | -0.207650 | 0.096307  | 0.051588  | ... | -0.005230 | -0.291977 | 0.028039  | 0.037419  |   |
| 1 | 0.026485  | -0.012477 | 0.066108  | ... | -0.178015 | -0.127297 | -0.054074 | -0.065085 |   |
| 2 | -0.658986 | -0.141501 | -0.186631 | ... | 0.053974  | 0.185881  | -0.007385 | 0.085401  |   |
| 3 | 0.032491  | 0.134994  | 0.117323  | ... | 0.083053  | 0.379816  | 0.122336  | -0.157672 |   |
| 4 | 0.013562  | 0.152696  | -0.084538 | ... | 0.270027  | -0.331905 | 0.051465  | 0.050213  |   |

|   | 24        | 25        | 26        | 27        | 28        | 29        |
|---|-----------|-----------|-----------|-----------|-----------|-----------|
| 0 | 0.238741  | -0.248865 | -0.007800 | 0.621952  | 0.029046  | -0.106701 |
| 1 | -0.223562 | -0.695037 | 0.198811  | -0.371491 | 0.046585  | 0.098040  |
| 2 | -0.021493 | -0.033900 | 0.006939  | -0.078475 | 0.070041  | 0.066776  |
| 3 | -0.304626 | 0.069615  | -0.078672 | -0.166916 | 0.101930  | -0.138801 |
| 4 | -0.373392 | -0.089393 | -0.085101 | 0.056272  | -0.084307 | -0.087949 |

[5 rows x 30 columns]

```
[37]: # Presenting the result
print("ETF Returns Matrix Dimension:")
print(daily_returns.shape)
print("\nDimension of Matrix U:")
print(U_st_return.shape)
print("\nSingular values:")
print(s_st_return**2)
print("\nDimension of Matrix V^T:")
print(VT_st_return.shape)
```

ETF Returns Matrix Dimension:  
(251, 30)

Dimension of Matrix U:  
(251, 251)

Singular values:

|             |            |            |            |            |             |
|-------------|------------|------------|------------|------------|-------------|
| 16.94028088 | 1.84467656 | 1.26660112 | 1.0033836  | 0.85284607 | 0.78526804  |
| 0.68313883  | 0.62148367 | 0.56747695 | 0.54817649 | 0.50067321 | 0.4355354   |
| 0.39437057  | 0.36516273 | 0.35232383 | 0.33001439 | 0.29732338 | 0.28207821  |
| 0.24393001  | 0.23422616 | 0.21981263 | 0.21796549 | 0.19353803 | 0.17443758  |
| 0.16274534  | 0.13109828 | 0.10123079 | 0.09744171 | 0.08276568 | 0.06999438] |

Dimension of Matrix V^T:  
(30, 30)

```
[38]: # Visualization for SVD
svd = TruncatedSVD(n_components=standardized_returns_cov.shape[1] - 1)
svd_components = svd.fit_transform(x)
svdDf = pd.DataFrame(data=svd_components, columns=['singular value ' + str(i)
    ↪for i in range(1, standardized_returns_cov.shape[1])],
    ↪index=standardized_returns_cov.index)

plt.figure(figsize=(10, 6))
plt.plot(svd.singular_values_)
plt.xlabel('Singular Value Index')
plt.ylabel('Singular Value')
plt.title('Singular Values from SVD')
plt.show()
```

